

TensorBoard

Glenn Bruns
CSUMB

Much material in this deck from Géron, Hands-on Machine Learning with Scikit-Learn and TensorFlow

Learning outcomes

After this lecture you should be able to:

1. Extend TensorFlow code to enable use of TensorBoard
2. Use TensorBoard
3. Use the following language features of TensorFlow:
 - saving and restoring models
 - name scopes
 - modular graph construction

Recap: TensorFlow

```
n_epochs = 1000
learning_rate = 0.01
```

```
X = tf.constant(scaled_housing_data_plus_bias, dtype=tf.float32, name="X")
y = tf.constant(housing.target.reshape(-1, 1), dtype=tf.float32, name="y")
theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0, seed=42),
name="theta")
y_pred = tf.matmul(X, theta, name="predictions")
error = y_pred - y
mse = tf.reduce_mean(tf.square(error), name="mse")
gradients = 2/m * tf.matmul(tf.transpose(X), error)
training_op = tf.assign(theta, theta - learning_rate * gradients)

init = tf.global_variables_initializer()
```

```
with tf.Session() as sess:
    sess.run(init)

    for epoch in range(n_epochs):
        if epoch % 100 == 0:
            print("Epoch", epoch, "MSE =", mse.eval())
            sess.run(training_op)

    best_theta = theta.eval()
```

Saving models

```
n_epochs = 1000
learning_rate = 0.01

X = tf.constant(scaled_housing_data_plus_bias, dtype=tf.float32, name="X")
...
gradients = 2/m * tf.matmul(tf.transpose(X), error)
training_op = tf.assign(theta, theta - learning_rate * gradients)
init = tf.global_variables_initializer()
saver = tf.train.Saver()
```

```
with tf.Session() as sess:
    sess.run(init)

    for epoch in range(n_epochs):
        if epoch % 100 == 0:          # checkpoint every 100 epochs
            save_path = saver.save(sess, "/tmp/my_model.ckpt")
            sess.run(training_op)

    best_theta = theta.eval()
    save_path = saver.save(sess, "/tmp/my_model_final.ckpt")
```

"saving a model" = saving values of all variables in the model

Restoring models

```
n_epochs = 1000
learning_rate = 0.01

X = tf.constant(scaled_housing_data_plus_bias, dtype=tf.float32, name="X")
...
gradients = 2/m * tf.matmul(tf.transpose(X), error)
training_op = tf.assign(theta, theta - learning_rate * gradients)
init = tf.global_variables_initializer()
saver = tf.train.Saver()

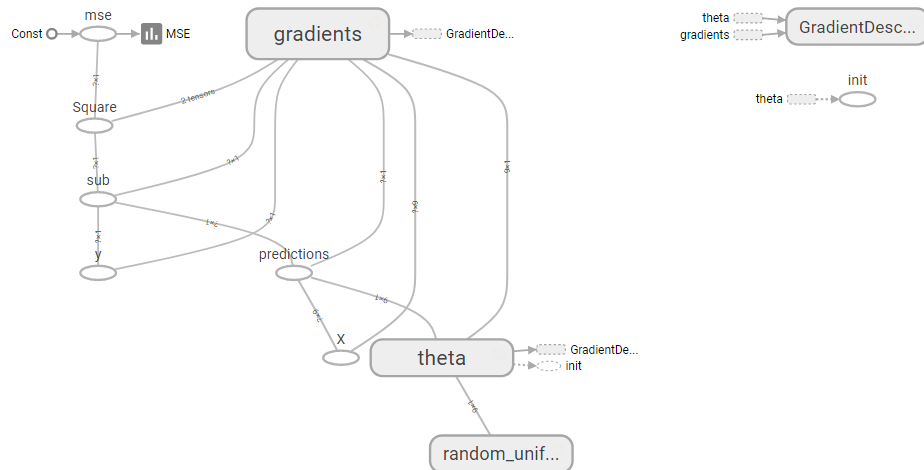
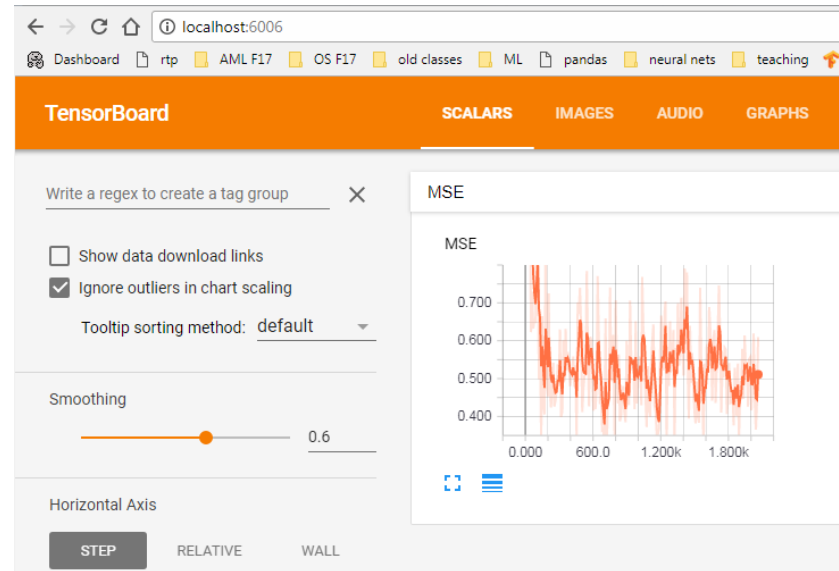
with tf.Session() as sess:
    # sess.run(init)
    saver.restore(sess, "/tmp/my_model_final.ckpt") # instead of run(init)
    for epoch in range(n_epochs):
        if epoch % 100 == 0: # checkpoint every 100 epochs
            save_path = saver.save(sess, "/tmp/my_model.ckpt")
            sess.run(training_op)
    best_theta = theta.eval()
```

By default, variables are saved and restored under their own names.

TensorBoard

Provided with TensorFlow; allows browser-based visualization to:

- Visualize data collected during graph execution
 - e.g., to see training progress
- Visualize the computation graph itself



Modifying code for TensorFlow

```
now = datetime.utcnow().strftime("%Y%m%d%H%M%S")
root_logdir = "E:/Glenn/tf_logs"
logdir = "{} /run-{}".format(root_logdir, now)
```

start of program:
define unique log
directory

```
mse_summary = tf.summary.scalar('MSE', mse)
file_writer = tf.summary.FileWriter(logdir,
                                    tf.get_default_graph())
```

end of construction
phase: create node
to evaluate MSE

```
[...]
for batch_index in range(n_batches):
    X_batch, y_batch = fetch_batch(epoch, batch_index, batch_size)
    if batch_index % 10 == 0:
        summary_str = mse_summary.eval(feed_dict={X: X_batch, y: y_batch})
        step = epoch * n_batches + batch_index
        file_writer.add_summary(summary_str, step)
    sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
[...]
```

execution phase:
write MSE values
every so often

```
file_writer.close()
```

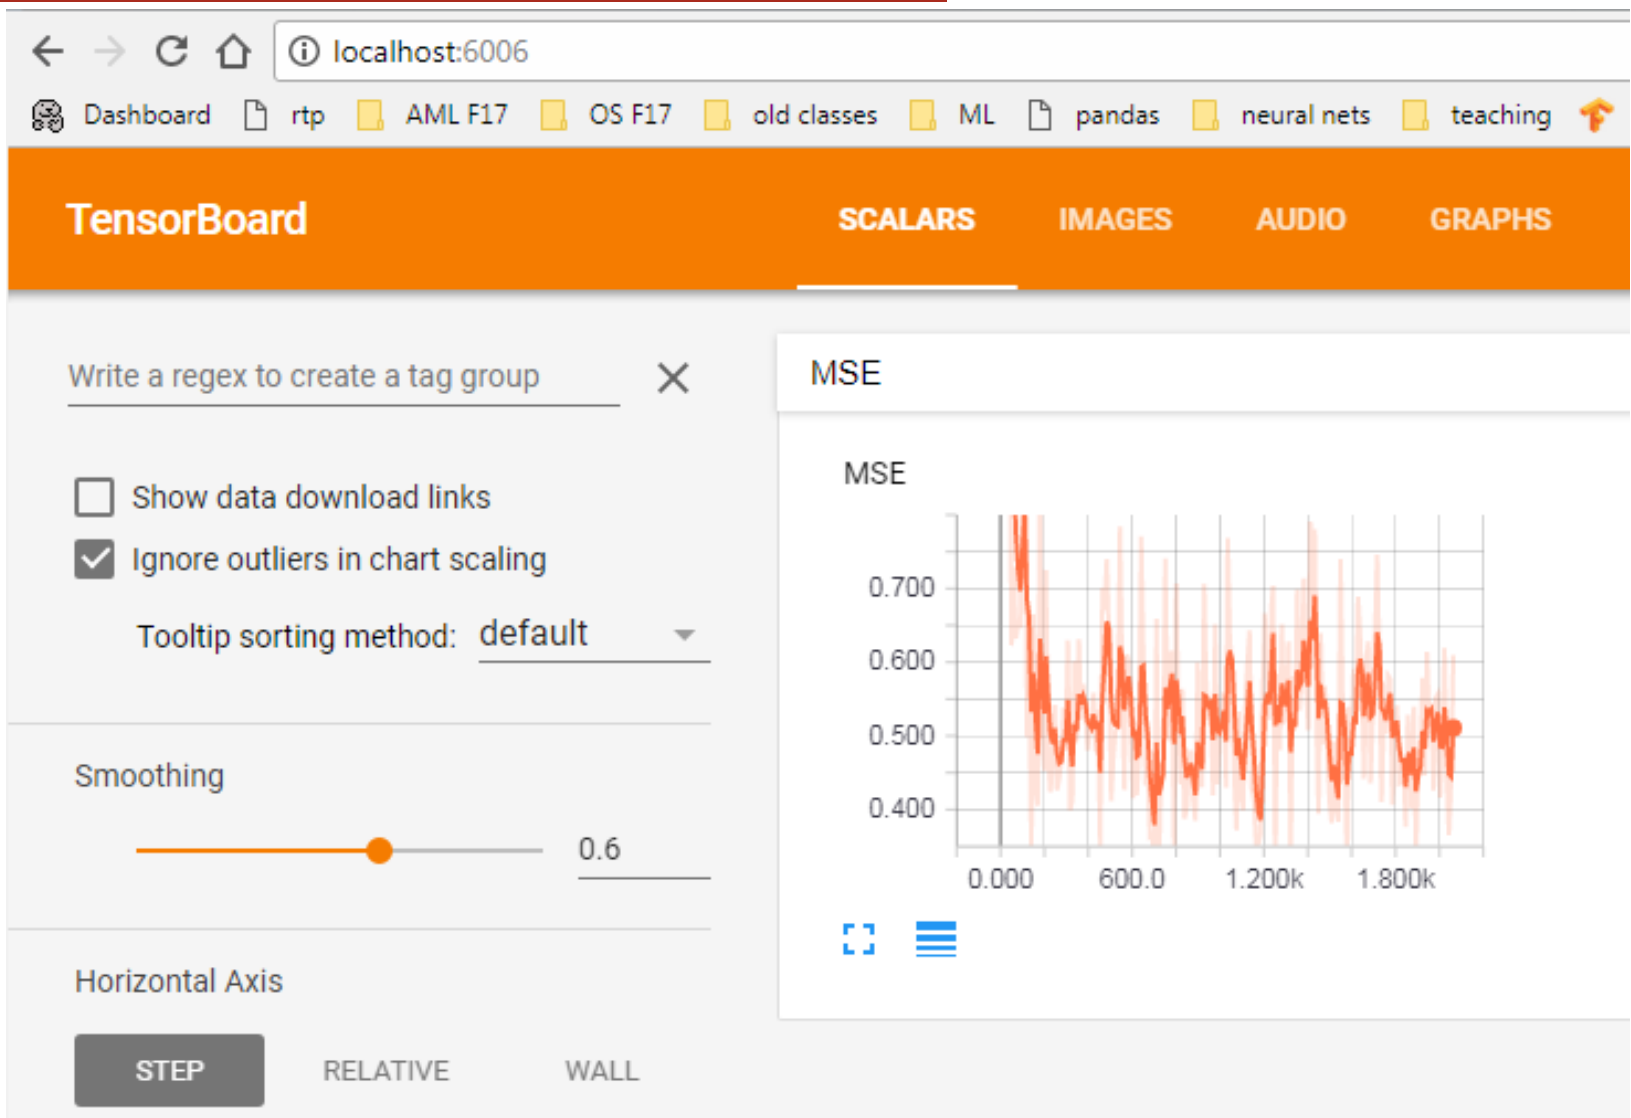
end of program:
close FileWriter

Running TensorBoard

After running your TF program and creating the log directory:

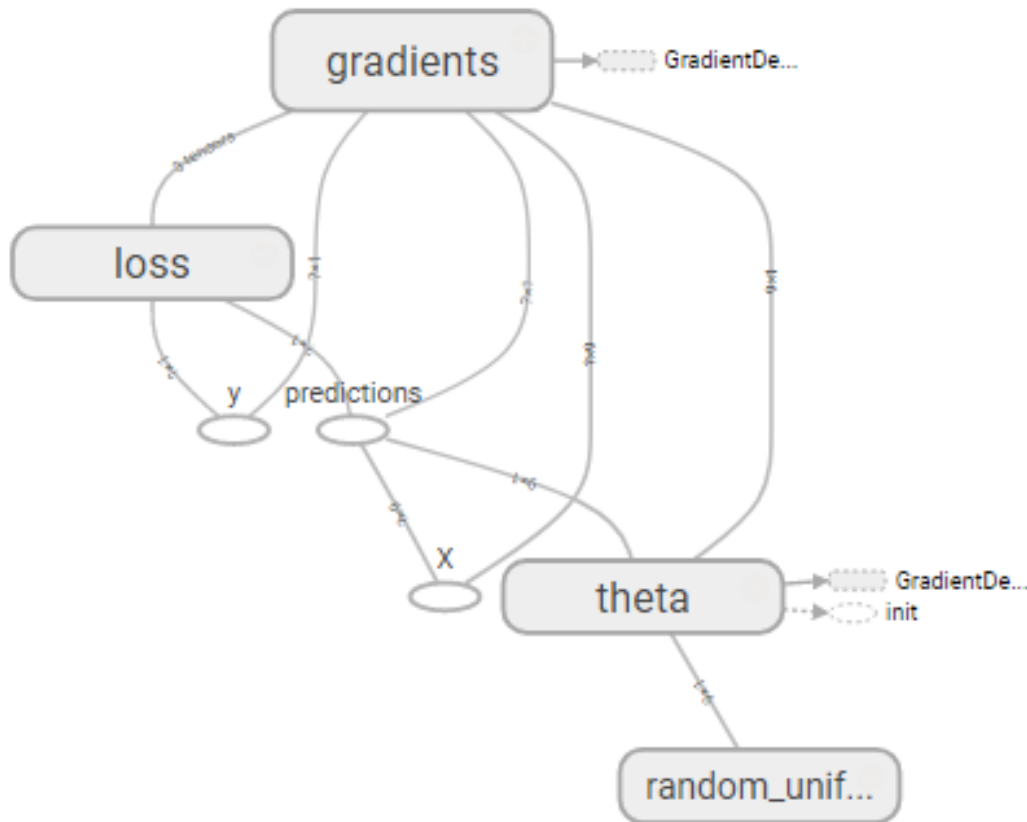
1. Within an Anaconda window, activate your tensorflow environment
2. `> tensorboard --logdir tf_logs/`
3. Enter URL localhost:6006 in your browser
 - or, try 0.0.0.0:6006 (this didn't work for me)
4. Click on MSE to see line plot of MSE values
5. Click on 'GRAPHS' tab to see your graph

MSE plot




Graph visualization

Main Graph



Auxiliary Nodes



- scroll to zoom in and out
- hold left mouse button to pan
- hover over node, then click + to expand a subgraph

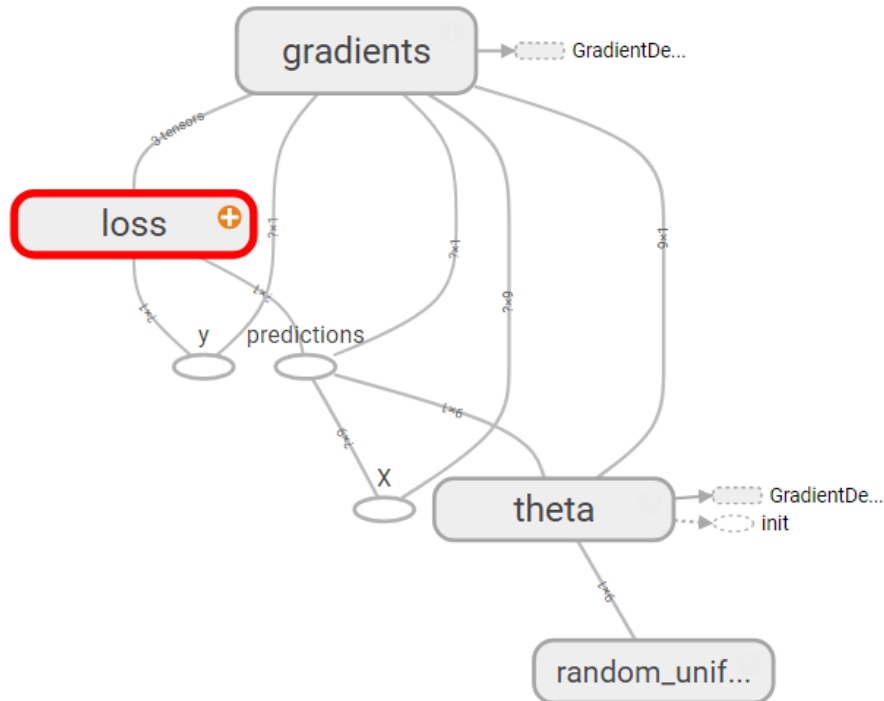
Name scopes

Reduce clutter in TensorBoard view of graph by grouping related nodes.


```
theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0, seed=42),  
name="theta")  
y_pred = tf.matmul(X, theta, name="predictions")  
  
# put error and mse within a single name scope  
with tf.name_scope("loss") as scope:  
    error = y_pred - y  
    mse = tf.reduce_mean(tf.square(error), name="mse")  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)  
training_op = optimizer.minimize(mse)
```

This code creates a "loss" name scope that simplifies the graph

Loss namespace



collapsed



expanded

Duplicated TF code

In any language, duplicated code leads to errors and maintenance problems. (DRY principle)

```
n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")

w1 = tf.Variable(tf.random_normal((n_features, 1)), name="weights1")
w2 = tf.Variable(tf.random_normal((n_features, 1)), name="weights2")
b1 = tf.Variable(0.0, name="bias1")
b2 = tf.Variable(0.0, name="bias2")

z1 = tf.add(tf.matmul(X, w1), b1, name="z1")
z2 = tf.add(tf.matmul(X, w2), b2, name="z2")

relu1 = tf.maximum(z1, 0., name="relu1")
relu2 = tf.maximum(z1, 0., name="relu2")
output = tf.add(relu1, relu2, name="output")
```

$$h_{\mathbf{w},b}(\mathbf{X}) = \max(\mathbf{X} \cdot \mathbf{w} + b, 0) \quad (\text{rectified linear unit, ReLU})$$

Modularity

```
# build a ReLU
def relu(X):
    w_shape = (int(X.get_shape()[1]), 1)
    w = tf.Variable(tf.random_normal(w_shape), name="weights")
    b = tf.Variable(0.0, name="bias")
    z = tf.add(tf.matmul(X, w), b, name="linear")
    return tf.maximum(z, 0., name="relu")

n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = [relu(X) for i in range(5)]
output = tf.add_n(relus, name="output")
```

When a node is created, TF ensures name uniqueness when adding the node to the graph.

First ReLU contains "weights", "biases", etc.

Second ReLU contains "weights_1", "biases_1", etc.

Modularity improved with name scope


```
# build a ReLU
def relu(X):
    with tf.name_scope("relu"):
        w_shape = (int(X.get_shape()[1]), 1)
        w = tf.Variable(tf.random_normal(w_shape), name="weights")
        b = tf.Variable(0.0, name="bias")
        z = tf.add(tf.matmul(X, w), b, name="linear")
        return tf.maximum(z, 0., name="relu")

n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = [relu(X) for i in range(5)]
output = tf.add_n(relus, name="output")
```

The graph in TensorBoard will now look much better.

Modularity + name scope in TB

Main Graph



As our computation graphs get bigger, the use of modularity and name scope will become more important.

Summary

- ❑ saving and restoring sessions
- ❑ TensorBoard
 - modifying your code for TensorBoard
 - using TensorBoard
- ❑ name scopes and modularity to clarify code and its TensorBoard graph