

# *Exam 2 preview*

---

Glenn Bruns  
CSUMB

much of this material is based on Geron's "Hand-on" text

# Coverage of exam

---

All material up to CNNs

Focus on material since the last exam:

- dimensionality reduction
- TensorFlow and neural networks

However, exam will also include:

- linear algebra, training models, support vector machines, ensemble learning

# Structure of exam

---

## 1. Concepts and theory (25 mins)

- conceptual questions
- questions about the math
- paper and pencil – no notes or other resources

## 2. Practical (45 mins)

- add your code to starter iPython notebook
- use any resources you like

# How to prepare

---

- Note **learning outcomes** at the front of each slide deck
  - ask yourself if you can do these things
- Practice on **lab and homework problems**
- Don't passively review lecture slides
  - **actively review** by writing test questions
  - make flash cards for yourself, and use them

# Topics since exam 1

---

- Dimensionality reduction
  - PCA (part of chap 8)
- Neural nets
  - TensorFlow (chap 9)
  - Neural nets (chap 10)
  - Training deep neural nets (chap 11)
  - CNNs (chap 13)

# Dimensionality reduction: PCA

---

- ❑ High-dimensional data in machine learning
- ❑ Issues with high-dimensional data
- ❑ Feature selection and dimensionality reduction
- ❑ Projection and manifold learning
- ❑ PCA concept
- ❑ Computing the principal components
- ❑ PCA with Numpy and with Scikit-Learn
- ❑ Deciding on number of components to project to

# Dimensionality reduction

---

Two main approaches:

- **projection** (PCA and its relatives)


- assumption: training data is not spread out uniformly across all dimensions – training instances lie in a much lower-dimensional subspace
- method: project data onto lower-dimensional subspace

- **manifold learning** (LLE, Isomap, MDS, t-SNE,...)

- a 2D manifold is a 2D shape that can be bent and twised in a higher-dimensional space
- assumption: most high-dimension datasets lie close to a much lower-dimension manifold
- method: model the manifold on which the data lies

# Principle component analysis: concept

When we are given 3-D data, the values are relative to a coordinate system, with x, y, z axes.



Think of this as expressing the data relative to three vectors, one for each axis:

$$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

We could equally well express our data relative to three other orthogonal axes.

Which other axes are best?



# Principal component analysis

Main idea:

- The single best axis is the one with most variance
- The next best axis is the one, of those orthogonal to the best axis, with the most variance




figure: [austingwalters.com/pca-principal-component-analysis/](http://austingwalters.com/pca-principal-component-analysis/)

# Computing the principal components

---

One method is to use Singular Value Decomposition:

$$X = U \Sigma V^T$$

Recall that if  $X$  is an  $m \times n$  matrix, then  $V^T$  is an  $n \times n$  orthogonal matrix.

The columns  $c_1, \dots, c_n$  of  $V$  are the principal components:

$$\mathbf{V} = \begin{pmatrix} | & | & & | \\ \mathbf{c}_1 & \mathbf{c}_2 & \cdots & \mathbf{c}_n \\ | & | & & | \end{pmatrix}$$

this is a corrected version of Fig. 8-1 in the Géron text

# Projecting down to d dimensions

$$X_d = X W_d$$

Here  $W_d$  is the first  $d$  columns of  $V$ . Note the sizes:

$$m \times n \quad n \times d \quad \rightarrow \quad m \times d$$

Example: suppose our data consists of 3 instances, and we want to reduce from 4 to 2 features.

$$\begin{array}{ccc} X & W_d & X_d \\ \begin{pmatrix} x_1 & x_2 & x_3 & x_4 \end{pmatrix} & \begin{pmatrix} a \\ b \\ c \\ d \end{pmatrix} & = \begin{pmatrix} ? \\ \end{pmatrix} \end{array}$$


an original feature vector

first principal component

# How many dimensions to project to?

---

```
pca = PCA()  
pca.fit(X)  
cumsum = np.cumsum(pca.explained_variance_ratio_)  
d = np.argmax(cumsum > 0.95) + 1  
  
# the easier way  
pca = PCA(n_components = 0.95)  
X_reduced = pca.fit_transform
```



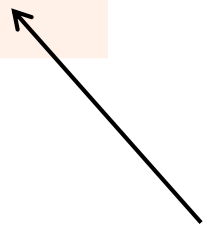
# Kernel PCA

---

PCA provides a linear transformation of the original data.

What if we need a non-linear transformation?

	SVM	PCA
linear case	linear SVM	PCA
non-linear case	SVM with polynomial or RBF kernel	kernel PCA



# Tuning kernel PCA, example

---

A nice example from Géron showing the use of a Pipeline as a predictor in GridSearchCV.

```
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline

clf = Pipeline([
    ("kpca", KernelPCA(n_components=2)),
    ("log_reg", LogisticRegression())
])

param_grid = [{
    "kpca__gamma": np.linspace(0.03, 0.05, 10),
    "kpca__kernel": ["rbf", "sigmoid"]
}]

grid_search = GridSearchCV(clf, param_grid, cv=3)
grid_search.fit(X, y)
```

# Computational complexity

---

PCA:  $O(m \times n^2) + O(n^3)$

Randomized PCA:  $O(m \times d^2) + O(d^3)$

LLE:  $O(m \log(m)) \log(k)$  (find k-nearest neighbors)  
       $+ O(mnk^3)$  (reconstruction cost)  
       $+ O(dm^2)$  (embedding cost)

## Recall:

$m$  number of instances

$n$  number of features

$k$  number of nearest neighbors

$d$  # of dimensions in lower-dimensional space

# Neural nets: introduction to TF

---

- ❑ computation model
- ❑ handling sessions
- ❑ evaluating a node
- ❑ linear regression example




# What is TensorFlow?

It's an open source library for numerical computation.

Principle:

1. define a **computation graph**
2. execute the graph efficiently with C++ code



You can even break the graph into pieces and run them on separate CPUs or GPUs.

# Neural nets: linear regression with TF

---

- linear regression as an optimization problem
- TF code for linear regression
  - computing gradients
  - selecting an optimizer
  - mini-batch

# Using TensorFlow's gradient descent

---

```
...
y_pred = tf.matmul(X, theta, name="predictions")
error = y_pred - y
mse = tf.reduce_mean(tf.square(error), name="mse")
gradients = tf.gradients(mse, [theta])[0]
training_op = tf.assign(theta, theta - learning_rate * gradients)
```

You can use one of TensorFlow's built-in optimizers instead:

```
...
y_pred = tf.matmul(X, theta, name="predictions")
error = y_pred - y
mse = tf.reduce_mean(tf.square(error), name="mse")
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
training_op = optimizer.minimize(mse)
```

Switching to an alternative TF optimizer is super-easy.

# Neural nets: intro

---

- ❑ artificial neuron
- ❑ perceptron
- ❑ linear threshold unit
- ❑ multi-layer perceptron
- ❑ activation function
- ❑ perceptron learning rule

# Artificial neurons

---


A first mathematical model of neurons, by McCulloch and Pitts.

**artificial neuron:**

- one or more binary inputs
- one binary output
- output activated when a certain number of inputs are active

**example:**


- one input, two outputs
- activate output when both inputs active



# Linear Threshold Unit LTU


An LTU is another kind of artificial neuron.

- ❑ It has numeric, not boolean, inputs and outputs.
- ❑ It computes a weighted sum of its inputs.
- ❑ The output is a step function applied to the weighted sum.



# Perceptrons

A perceptron is a single layer of LTUs, with each neuron connected to all the inputs.



The figure shows a perceptron with 2 inputs and 3 outputs. The "input neurons" output whatever input they are fed.


# Multi-Layer Perceptron

A multi-layer perceptron has:

- one input layer
- one or more layers of LTUs (hidden layers)
- one final layer of LTUs (output layer)

Every layer but the output layer:

- includes a bias neuron
- is fully connected to the next layer





# Activation functions

Key to backprop:

- instead of step function, use **logistic function**:

$$\sigma(\mathbf{z}) = \frac{1}{1 + \exp(-z)}$$



This is called an **activation function**. Other popular activation functions:

- **hyperbolic tangent function**:

$$\tanh(\mathbf{z}) = 2\sigma(2z) - 1$$

- **ReLU function**:

$$\text{ReLU}(\mathbf{z}) = \max(0, z)$$




# Neural nets: backpropagation

---

- gradient descent
- gradient descent in a computation graph
- backpropagation algorithm

# Backprop: example with two nodes



$$z = y^2$$

$$y = x + x^3$$

Approach 1: combine  $f$  and  $g$  to get  $z = (x + x^3)^2$  and then do the same thing as the last example.


Approach 2: get  $\frac{dz}{dx}$  by using the chain rule:  $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$

Example: Let the input  $x$  be 1 (so  $y$  is 2 and  $z$  is 4). How to adjust  $x$  to make  $z$  smaller?

Work out that  $\frac{dz}{dy}$  is  $2y$ , and  $\frac{dy}{dx}$  is  $1 + 3x^2$ . Then the chain rule says that the slope of the  $f, g$  combined, at  $x=1$ , is:

$$\frac{dz}{dy} (2) * \frac{dy}{dx} (1) = 4 * 4 = 16. \text{ For new } x, \text{ use } x - \eta \frac{dz}{dx}(x)$$

# A weighted sum node



$$z = y^2$$

$$y = w_1 x_1 + w_2 x_2$$

Node  $\Sigma$  outputs the weighted sum of its inputs. The coefficient values are  $w_1 = 0.5$  and  $w_2 = 2.0$ .

How to modify **the coefficients of  $\Sigma$**  to make  $z$  smaller? (Now treat  $x_1$  and  $x_2$  as constants.)

Use the multi-variable chain rule:  $\frac{\partial z}{\partial w_1} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial w_1}$  (similarly for  $w_2$ )

Example: Let  $x_1 = 1$  and  $x_2 = 3$  (so  $y = 0.5 + 6 = 6.5$ ).

1.  $\frac{\partial z}{\partial y}$  is  $2y$ , and  $\frac{\delta y}{\delta w_1}$  is  $x_1$  (note:  $\frac{\delta y}{\delta w_1}$  does not depend on  $w_1$  !)
2. By chain rule:  $\frac{\partial z}{\partial w_1}$  at  $w_1, w_2$  is  $\frac{dz}{dy}(6.5) * \frac{dy}{dw_1}(0.5, 2.0) = 13 * 1 = 13$
3. For new value of  $w_1$ , use  $w_1 - \eta \frac{\delta z}{\delta w_1}(0.5, 2.0)$

# Notes from Goodfellow et al

---

- Multilayer perceptrons (MLPs) are also called **deep feedforward networks**
- **back-propagation** is not gradient descent – it refers only to the method for computing the gradient
- back-propagation is a very general technique:
  - is not limited to the gradient of a cost function with respect to its parameters
  - not limited to neural networks

source: Deep Learning, by Goodfellow, Bengio, and Courville

# Neural nets: training

---

- ❑ list some of the APIs built on top of TensorFlow
- ❑ build and train a DNN using TensorFlow
- ❑ tune the hyperparameters of your neural net

# Summary of building the net

---

1. define inputs
2. specify placeholders for training, target data
3. create the network
4. define a loss function
5. specify an optimizer
6. specify a performance measure
7. make initializer and saver

## Questions:

- does the optimizer depend on the loss function?
- does the performance measure depend on the optimizer?

# Tuning the network parameters

---

- number of hidden layers
  - with more layers, exponentially fewer total neurons can be used to model complex functions
  - hierarchical concept
  - start with a couple, ramp up
- number of neurons per hidden layers
  - fewer and fewer neurons per layer, as you move to output layer
- activation functions
  - often, ReLU for hidden layers (good performance)
  - output layer: often softmax for classification, nothing for regression

please read this section of our text carefully for details



# Training DNNs: vanishing gradients

---

- diagnose and address the vanishing and exploding gradients problem

# Causes of the problem

---

"Understanding the Difficulty of Training Deep Feedforward Neural Networks" (Glorot and Bengio, 2010), diagnosed the issue.

1. Use of the sigmoid activation function
2. Method used to initialize network parameters
  - random initialization with mean 0, std dev 1

Idea:

- with 1 and 2 above, the outputs of a layer have greater variance than the inputs of a layer
- activation function then saturates in top layers

# Idea 2: dump the sigmoid activation


## Replacement candidate 1: ReLU activation function

pros:

- doesn't saturate for positive values
- fast to compute

cons:

- dying ReLU problem:  
many neurons in network  
stop outputting anything  
other than 0
- if weighted sum of  
neuron's inputs is  
negative, ReLU gives 0
- once this happens, it  
tends to stay that way



# Dump the sigmoid (cont'd.)


## Replacement candidate 2: leaky ReLU

pros:

- unlike a ReLU, it can never "die"
- a recent paper claims it always outperforms plain ReLU

## Other cousin candidates:

- randomized leaky ReLU
  - $\alpha$  is picked randomly during training
- parametric leaky ReLU
  - $\alpha$  is learned during training



$\alpha$  is a hyperparameter;  
typically set to 0.01


values of up to 0.2  
may work even better

# Dump the sigmoid (cont'd.)

## Replacement candidate 3: exponential linear unit (ELU)

pros:

- outperformed all ReLU variants in a recent study
- training time reduced; performance better
- function is smooth everywhere, helping with gradient descent
- negative value when  $z < 0$ , avoids vanishing gradients
- non-zero gradient when  $z < 0$ , avoids dying units



$$\text{ELU}_{\alpha}(z) = \begin{cases} \alpha(\exp(z) - 1) & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases}$$

hyperparameter usually set to 1

cons: slower to compute than ReLU and its variants

# Solution idea 3: Batch normalization

---

He initialization plus better activation function especially helps with vanishing/exploding gradients during beginning of training.

A newer technique (2015) is **batch normalization**.

Idea:

- ❑ add operation just before activation function
- ❑ this operation lets the model learn the optimal scale and mean of the inputs for each layer
- ❑ inventors claim many benefits, including reduction in vanishing gradients, and better performance

See textbook for using batch normalization in TensorFlow

# Training DNNs: optimizers

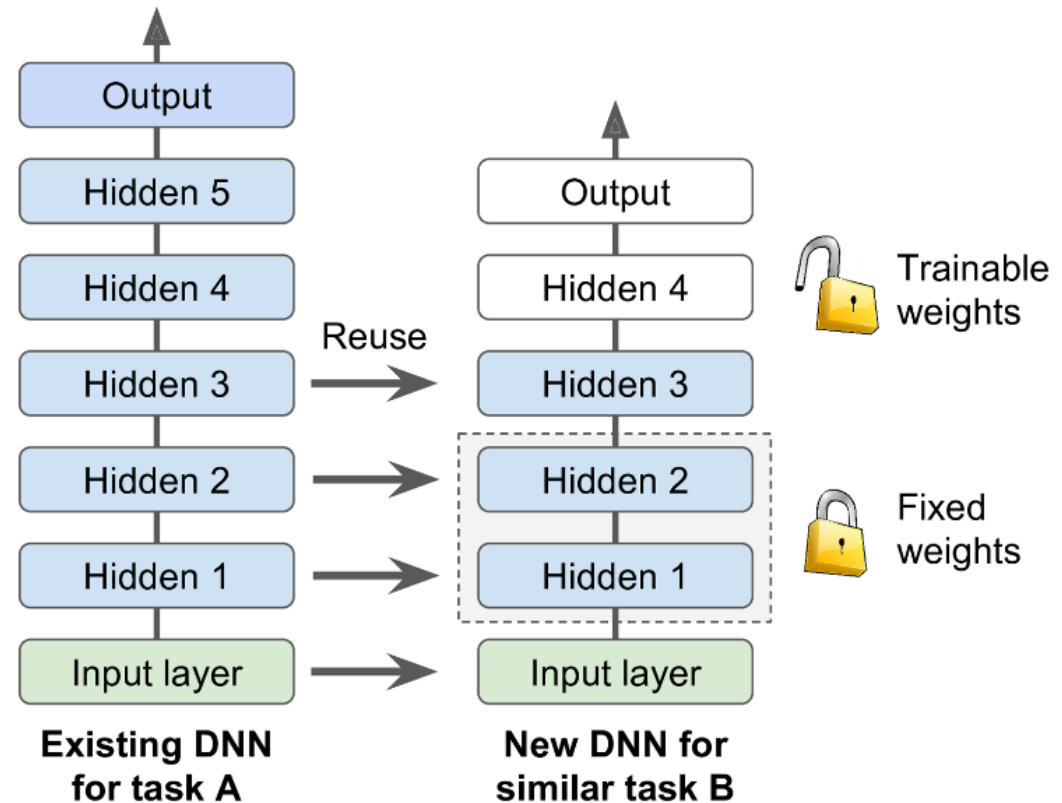
---

- ❑ reuse parts of TensorFlow models
- ❑ select a fast optimizer

# Reusing layers

- Training large nets takes a long time
- Many nets solve similar tasks

**Transfer learning:** reuse lower levels of a net solving a problem similar to your own



Example: you want to classify vehicles from images; existing nets classify images of animals, vehicles, etc.



# Caching layers

---

If the reused layers are frozen, you only need to compute the output of the topmost frozen layer once.

Strategy:

- ❑ run the whole training set through the lower levels
- ❑ don't batch training data – batch outputs from the topmost frozen layer
- ❑ use these during training

# Strategies for tweaking reuse

---

- we know we don't want to reuse the output layer of original model
- how many of the hidden layers should be reused?
  - the higher the level, the less likely to be used
- strategy:
  - start by freezing all borrowed layers
  - train your model, see how it performs
  - if performance not very good, unfreeze one or two top borrowed layers
  - if performance still not very good, drop or replace the topmost borrowed layers

# Momentum optimization

---

Idea: gain momentum on downhill slopes and "roll past" local minima

$$\mathbf{m} = \beta \mathbf{m} + \eta \nabla_{\theta} J(\theta)$$
$$\theta = \theta - \mathbf{m}$$

where  $\mathbf{m}$  is the **momentum vector**.

$0 \leq \beta \leq 1$     0 means "high friction", 1 means "no friction"

Suppose gradient is constant,  $\beta = 0.9$  (typical value)

Example:

$$\eta = 0.1$$

$$\nabla_{\theta} J(\theta) = (1, 1)$$

$$\mathbf{m} = (0, 0)$$

$$\theta = (0.5, 2)$$

$$\mathbf{m} = 0 + (0.1, 0.1)$$

$$\theta = (0.4, 1.9)$$

$$\mathbf{m} = ?$$

$$\theta = ?$$

$$\theta = (0.5, 2)$$

$$\mathbf{m} = 0 + (0.1, 0.1)$$

$$\theta = (0.4, 1.9)$$

$$\mathbf{m} = (0.09 + 0.09) + (0.1, 0.1)$$
$$= (0.19, 0.19)$$

$$\theta = (0.21, 1.71)$$

# Nesterov Accelerated Gradient

---

- A variant of Momentum optimization
- Idea: measure gradient at position slightly ahead of direction of momentum

$$\mathbf{m} = \beta \mathbf{m} + \eta \nabla_{\theta} J(\theta + \beta \mathbf{m})$$

$$\theta = \theta - \mathbf{m}$$

- In plain Momentum optimization we had

$$\mathbf{m} = \beta \mathbf{m} + \eta \nabla_{\theta} J(\theta)$$

"look where you're going"

# Training DNNs: learning rate, regl'zn

---


- ❑ learning rate scheduling in DNNs
- ❑ regularization in DNNs

# Learning rate scheduling

---

## Predetermined piecewise constant learning rate:

- reduce learning rate every so many epochs
- example: initially  $\eta = 0.1$ , then 0.001 after 50 epochs



## Performance scheduling:



- measure test error every N steps
- reduce learning rate by a factor of  $\lambda$  when the error stops dropping

# Exponential scheduling

- learning rate a function of iteration number:

$$\eta(t) = \eta_0 10^{-t/r}$$

- two tuning parameters:  $\eta_0, r$




# Early stopping

---

Stop training when performance on test set starts dropping

One way, in TensorFlow:

- ❑ periodically evaluate model on test set
- ❑ save a 'winner' snapshot if it outperforms earlier winner snapshots
- ❑ stop training if number of steps since last winner exceeds a threshold (like 2000 steps)





# Dropout

Super effective and popular.

Idea is simple:

- at each training step, each neuron (except output neurons) has some probability  $p$  of being ignored
- dropout rate  $p$  is typically set to 0.5
- dropout only occurs during training




figure source: Géron

# Practical guidelines

---

Configure your DNN like this by default:

- initialization: He initialization
- activation function: ELU
- normalization: batch normalization
- regularization: dropout
- optimizer: Adam
- learning rate schedule: none

These guidelines are straight out of Géron's book.

# Practical guidelines, cont'd.

---

## Tweaking your configuration:

- ❑ Can't find good learning rate? Try a learning schedule, such as exponential decay
- ❑ Training set too small? Try data augmentation
- ❑ Need a sparse model? Add  $l_1$  regularization, or FTRL instead of Adam optimization
- ❑ Need a fast model at runtime? Drop batch normalization, replace ELU with leaky ReLU

# Convolutional neural nets


---

- explain the biological motivation for convolutional neural nets
- define the concepts of:
  - convolutional layer
  - padding
  - stride
  - filter
  - feature map
- write the expression defining the output of a neuron in a convolutional layer

# Convolutional layer


---

- in the first **convolutional layer** of a Convolutional Neural Network, each neuron "sees" only pixels in its **receptive field**
- in the second convolutional layer, each neuron is connected only to neurons in a small rectangle of the previous layer



# Padding

How to deal with this situation?




Idea 1: use a second layer that is smaller than the first

Idea 2: add padding around the edge of the first layer

"zero padding"

# Stride


A receptive field can skip over rows and columns of a layer



# Filters

---

The weights associated with neuron on the right can be shown as an image – like a heatmap







# Feature maps

---

A layer full of neurons using the same filter is a **feature map**.



# Stacking feature maps



a convolutional layer applies multiple filters to its inputs

a convolutional layer can then detect multiple features anywhere in its inputs

# Pooling

---

Basically: aggregate over receptive field

As with convolution, you must specify:

- size of the rectangle
- stride
- padding type

But now, instead of weights, also specify:


- aggregation function (max, mean, etc.)

Reduces computational load, memory usage, and number of parameters


Question: how is overfitting affected?

# Pooling

---




Convolved  
feature



Pooled  
feature

# Max pooling example




Notes:

- pooling is usually done on each channel independently
- but, pooling can also be done across channels

# CNN architecture

---

How are convolution and pooling combined in common CNNs?



- ❑ The convolutional layers usually have ReLU activation
- ❑ Image gets smaller but also deeper (more feature maps) as it goes through network
- ❑ Final layer outputs prediction (e.g. softmax layer)